

Decision Tree

- The decision tree from the name itself signifies that it is used for making decisions from the given dataset.
- The concept behind the decision tree is that it helps to select appropriate features for splitting the tree into subparts.
- the algorithm used behind the splitting is ID3.

For checking the purity, we have 2 features:

1. **Entropy:** It builds an appropriate decision tree for selecting the best splitter.

$$H(s) = -P_{(+)} \log_2 P_{(+)} - P_{(-)} \log_2 P_{(-)}$$

Here $P_{(+)}$ / $P_{(-)}$ = % of + ve class 1% of - ve class

Entropy can be defined as a measure of the purity of the sub split. The value lies between 0 and 1.

The algorithm calculates the entropy of each feature using the above formula, after every split and as the splitting continues, it selects the best feature and starts splitting according to it.

2. **Gini Impurity:** Gini impurity is a metric to measure how often a randomly chosen element would be incorrectly identified.


$$GiniIndex = 1 - \sum_j p_j^2$$

It means an attribute with lower Gini index should be preferred.

In the words of Wikipedia, its main goal is to measure how often a randomly chosen element from the set would be incorrectly labeled.

Difference between Entropy and Gini Indexing:

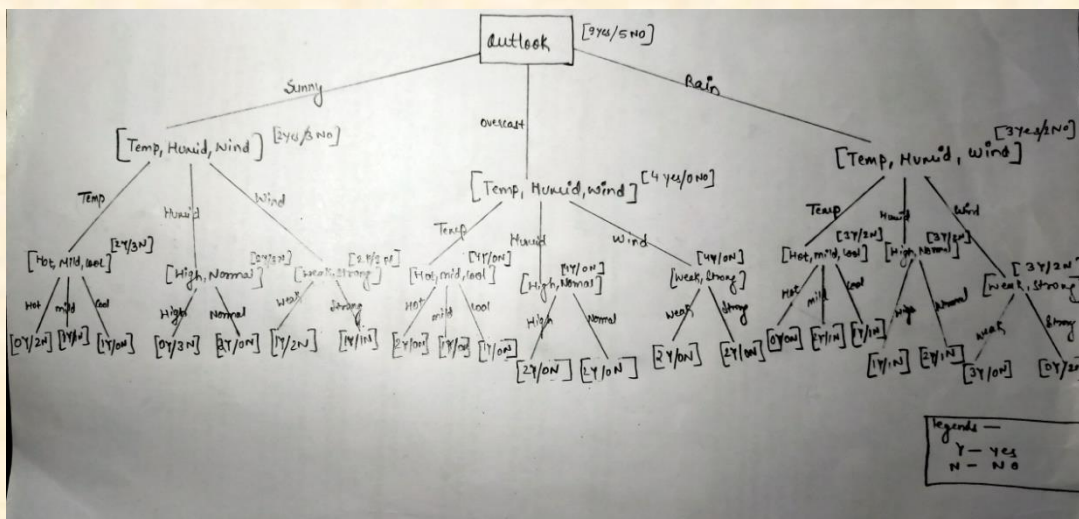
Entropy	Gini Indexing
1. Entropy can be defined as a measure of the purity of the sub split	1. It shows that the data point that we have randomly chosen for the splitting from the dataset how it is incorrectly labeled.
2. It is used for C4.5 decision tree.	2. It is used for CART
3. Higher values give inappropriate result,	3. It gives more accurate value and less than the entropy index which is its best quality, the small values show less impurity.
4. Takes more time	4. Takes less time
5. Entropy has a maximum impurity of 1 and maximum purity is 0.	5. The Gini index has a maximum impurity is 0.5 and maximum purity is 0



Playing Tennis Example

Day	Outlook	Temp	Humid	Wind	Play?
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Weak	Yes
4	Rain	Mild	High	Weak	Yes
5	Rain	Cool	Normal	Weak	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
10	Rain	Mild	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Weak	Yes
14	Rain	Mild	High	Strong	No

Decision tree w.r.t above Tennis dataset



Link for clear pic [image](#)

Practical Implementation

In [3]:

```
import pandas as pd
import numpy as np
```

In [4]:

```
df=pd.read_csv("https://raw.githubusercontent.com/shrikant-temburwar/Wine-Quality-Dataset/master/winequality-red.csv", sep=';')
```

In [5]:

df

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

1599 rows x 12 columns

In [6]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates             1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

In [7]:

df.isnull().sum()

Out[7]:

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density           0
pH                0
sulphates         0
alcohol           0
quality           0
dtype: int64
```

In [8]:

```
df.duplicated().sum()
```

Out[8]:

```
240
```

In [9]:

```
df.shape
```

Out[9]:

```
(1599, 12)
```

In [10]:

```
df=df.drop_duplicates()
```

In [11]:

```
1599-240
```

Out[11]:

```
1359
```

In [12]:

```
df.shape
```

Out[12]:

```
(1359, 12)
```

In [13]:

```
df['quality'].unique()
```

Out[13]:

```
array([5, 6, 7, 4, 8, 3])
```

In [14]:

```
df['quality'].value_counts()
```

Out[14]:

```
5    577
6    535
7    167
4     53
8     17
3     10
Name: quality, dtype: int64
```

In [15]:

```
from sklearn.model_selection import train_test_split, GridSearchCV
```

In [16]:

```
X=df.drop('quality',axis=1)
```

In [17]:

```
y=df['quality']
```

In [18]:

```
X
```

Out[18]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
5	7.4	0.660	0.00	1.8	0.075	13.0	40.0	0.99780	3.51	0.56	9.4
...
1593	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82	9.5
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0

1359 rows × 11 columns

In [19]:

```
y
```

Out[19]:

```
0      5
1      5
2      5
3      6
5      5
..
1593   6
1594   5
1595   6
1597   5
1598   6
Name: quality, Length: 1359, dtype: int64
```

In [20]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=10)
```

In [21]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [22]:

```
model=DecisionTreeClassifier()
```

In [23]:

```
model.fit(X_train,y_train)
```

Out[23]:

```
▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

In [24]:

```
model.score(X_train,y_train)
```

Out[24]:

1.0

In [25]:

```
y_pred=model.predict(X_test)
```

In [26]:

```
from sklearn.metrics import accuracy_score
```

In [27]:

```
accuracy_score(y_test,y_pred)
```

Out[27]:

0.5033407572383074

In [28]:

```
# Logistic Regression, SVM, ( RF, XGBoost, GB, AB)
```

In [35]:

```
grid_param={  
    'criterion':['gini','entropy'],  
    'max_depth':range(2,32,1),  
    'min_samples_leaf':range(1,10,1),  
    'min_samples_split':range(2,10,1),  
    'splitter':['best',"random"]}
```

In [36]:

```
from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=model,param_grid=grid_param,cv=5)
```

In [43]:

```
grid_search.fit(X_train,y_train)
```

Out[43]:

```
┌─── GridSearchCV ───┐  
│ ▶ estimator: DecisionTreeClassifier │  
│ ▶ DecisionTreeClassifier │  
└──────────────────┘
```

In [39]:

```
grid_search.best_params_  
# best parameter out of all
```

Out[39]:


```
{'criterion': 'gini',  
  'max_depth': 4,  
  'min_samples_leaf': 1,  
  'min_samples_split': 4,  
  'splitter': 'random'}
```

In [41]:

```
model_with_best_params=DecisionTreeClassifier(criterion= 'gini',  
  max_depth= 4,  
  min_samples_leaf= 1,  
  min_samples_split= 4,  
  splitter= 'random')
```

In [44]:

```
model_with_best_params.fit(X_train,y_train)
```

Out[44]:

```
▼ DecisionTreeClassifier  
DecisionTreeClassifier(max_depth=4, min_samples_split=4, splitter='random')
```

In [45]:

```
y_pred_2=model_with_best_params.predict(X_test)
```

In [46]:

```
accuracy_score(y_test,y_pred_2)
```

Out[46]:

```
0.5278396436525612
```

In [47]:

```
# Accuracy Increased
```

In []: